

1 Start coding or [generate](#) with AI.

1(a) Explanation : For this exercise, I will use the Boston Housing Dataset (for regression) and the Breast Cancer Wisconsin Dataset (for logistic regression classification). Since part (c) explicitly mentions logistic regression, I will focus on the Breast Cancer dataset for parts (a)-(g), which is a binary classification problem.

```
1 #1(a)
2 # Description: Load the Breast Cancer dataset (binary classification) for logistic regression.
3 # X will be the feature matrix (n_samples x n_features), y the binary target vector.
4
5 from sklearn.datasets import load_breast_cancer
6 import numpy as np
7
8 # Load dataset
9 data = load_breast_cancer()
10 X = data.data # Shape: (569, 30)
11 y = data.target # Shape: (569,) - 0: malignant, 1: benign
12
13 print(f"Dataset loaded: X shape {X.shape}, y shape {y.shape}")
```

Dataset loaded: X shape (569, 30), y shape (569,)

```
1 #1(b)
2 # Description: Randomly split data into 80% for training/validation and 20% for testing.
3 # We'll use sklearn's train_test_split.
4
5 from sklearn.model_selection import train_test_split
6
7 # Split 80/20
8 X_train_val, X_test, y_train_val, y_test = train_test_split(
9     X, y, test_size=0.2, random_state=42, stratify=y
10 )
11
12 print(f"Train+Val: {X_train_val.shape[0]} samples, Test: {X_test.shape[0]} samples")
```

Train+Val: 455 samples, Test: 114 samples

1(c) Objective of Logistic Regression with L2 Regularization

The objective function to minimize is:

$$J(w) = - (1/N) * \sum [ y^{(i)} * \log(\sigma(w^T x^{(i)})) + (1 - y^{(i)}) * \log(1 - \sigma(w^T x^{(i)})) ] + (\lambda/2) * ||w||_2^2$$

Where:

- N = number of training samples
- $\sigma(z) = 1 / (1 + \exp(-z))$  is the sigmoid function
- $\lambda$  = regularization strength (controls overfitting)
- $||w||_2^2$  = sum of squares of weights (L2 penalty)


This is the average cross-entropy loss plus L2 regularization.

```
1 #1(d)
2
3 # Description: Train logistic regression models with varying  $\lambda$  (L2 penalty).
4 # Plot:
5 # 1. Train/Test Cross-Entropy vs  $\log(\lambda)$ 
6 # 2. L2 norm of weights vs  $\log(\lambda)$ 
7 # 3. Individual weights vs  $\log(\lambda)$ 
8 # 4. Train/Test Accuracy vs  $\log(\lambda)$ 
9
10 from sklearn.linear_model import LogisticRegression
11 from sklearn.metrics import log_loss, accuracy_score
12 import matplotlib.pyplot as plt
13
14 # Define  $\lambda$  values
15 lambdas = [0, 0.1, 1, 10, 100, 1000]
16 C_values = [1e10 if lam == 0 else 1.0/lam for lam in lambdas] # sklearn uses C = 1/ $\lambda$ 
17
18 train_losses = []
19 test_losses = []
20 train_accs = []
21 test_accs = []
22 weight_norms = []
23 all_weights = [] # to store weight vectors
24
25 for lam, C in zip(lambdas, C_values):
```

```

26 # Initialize model
27 model = LogisticRegression(penalty='l2', C=C, solver='lbfgs', max_iter=1000, random_state=42)
28 model.fit(X_train_val, y_train_val)
29
30 # Predict probabilities for loss
31 y_train_proba = model.predict_proba(X_train_val)
32 y_test_proba = model.predict_proba(X_test)
33
34 # Compute cross-entropy (log loss)
35 train_loss = log_loss(y_train_val, y_train_proba)
36 test_loss = log_loss(y_test, y_test_proba)
37
38 # Compute accuracy
39 y_train_pred = model.predict(X_train_val)
40 y_test_pred = model.predict(X_test)
41 train_acc = accuracy_score(y_train_val, y_train_pred)
42 test_acc = accuracy_score(y_test, y_test_pred)
43
44 # Store weights and norm
45 w = model.coef_[0] # weights (bias is model.intercept_)
46 weight_norm = np.linalg.norm(w)
47
48 # Append to lists
49 train_losses.append(train_loss)
50 test_losses.append(test_loss)
51 train_accs.append(train_acc)
52 test_accs.append(test_acc)
53 weight_norms.append(weight_norm)
54 all_weights.append(w.copy())
55
56 # Convert to numpy for plotting
57 all_weights = np.array(all_weights) # shape: (6, 30)
58
59 # Plot 1: Cross-Entropy Loss
60 plt.figure(figsize=(12, 8))
61 plt.subplot(2, 2, 1)
62 plt.semilogx(lambdas, train_losses, 'o-', label='Train Loss')
63 plt.semilogx(lambdas, test_losses, 's-', label='Test Loss')
64 plt.xlabel('λ (log scale)')
65 plt.ylabel('Average Cross-Entropy')
66 plt.title('Train/Test Loss vs λ')
67 plt.legend()
68 plt.grid(True)
69
70 # Plot 2: L2 Norm of Weights
71 plt.subplot(2, 2, 2)
72 plt.semilogx(lambdas, weight_norms, 'o-', color='purple')
73 plt.xlabel('λ (log scale)')
74 plt.ylabel('||w||2')
75 plt.title('L2 Norm of Weights vs λ')
76 plt.grid(True)
77
78 # Plot 3: Individual Weights
79 plt.subplot(2, 2, 3)
80 for i in range(all_weights.shape[1]):
81     plt.semilogx(lambdas, all_weights[:, i], alpha=0.7, linewidth=0.8)
82 plt.xlabel('λ (log scale)')
83 plt.ylabel('Weight Value')
84 plt.title('Individual Weights vs λ')
85 plt.grid(True)
86
87 # Plot 4: Accuracy
88 plt.subplot(2, 2, 4)
89 plt.semilogx(lambdas, train_accs, 'o-', label='Train Acc')
90 plt.semilogx(lambdas, test_accs, 's-', label='Test Acc')
91 plt.xlabel('λ (log scale)')
92 plt.ylabel('Accuracy')
93 plt.title('Train/Test Accuracy vs λ')
94 plt.legend()
95 plt.grid(True)
96
97 plt.tight_layout()
98 plt.show()
99
100 """
101 ### Explanation of Results:
102 - As λ increases (stronger regularization), training loss ↑ (underfitting), test loss first ↓ then ↑ (optimal around λ=1~10).
103 - Weight norms shrink → regularization works.
104 - Individual weights are pulled toward zero.
105 - Test accuracy peaks at moderate λ, then drops → classic bias-variance trade-off.
106 - Scaling + higher max_iter fixed convergence warnings → results are now reliable.
107 """

```

 /usr/local/lib/python3.12/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT).

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT).

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT).

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT).


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```



1 #1(e)

2 # Description: Transform each feature using 5 Gaussian RBFs per feature.

3 # Means: evenly spaced from -10 to 10.  $\sigma \in \{0.1, 0.5, 1, 5, 10\}$

4

```
5 def gaussian_rbf(x, mu, sigma):
```

```
6     return np.exp(-0.5 * ((x - mu) / sigma) ** 2)
```

```
7
```

```
8 def apply_rbf_features(X, sigmas, n_centers=5):
```

```
9     """
```

```
10     For each feature in X, apply RBFs with given sigmas and fixed centers.
```

```
11     Returns expanded feature matrix.
```

```
12     """
```

```
13     centers = np.linspace(-10, 10, n_centers) # 5 centers
```

```
14     X_new_list = []
```

```
15
```

```
16     for sigma in sigmas:
```

```
17         X_sigma = []
```

```
18         for col in range(X.shape[1]): # for each original feature
```

```
19             x_col = X[:, col]
```

```
20             for mu in centers:
```

```
21                 rbf_val = gaussian_rbf(x_col, mu, sigma)
```

```
22                 X_sigma.append(rbf_val)
```

```
23             X_sigma = np.column_stack(X_sigma) # shape (n_samples, n_features * n_centers)
```

```
24             X_new_list.append(X_sigma)
```

```
25
```

```
26     # Concatenate all sigma-expanded features
```

```

27     X_new = np.hstack(X_new_list)
28     return X_new
29
30 # Define sigmas
31 sigmas = [0.1, 0.5, 1, 5, 10]
32
33 # Apply RBF transformation to train+val and test sets
34 X_train_val_rbf = apply_rbf_features(X_train_val, sigmas)
35 X_test_rbf = apply_rbf_features(X_test, sigmas)
36
37 print(f"Original feature dim: {X_train_val.shape[1]}")
38 print(f"RBF-expanded feature dim: {X_train_val_rbf.shape[1]}") # 30 features * 5 centers * 5 sigmas = 750

```

```

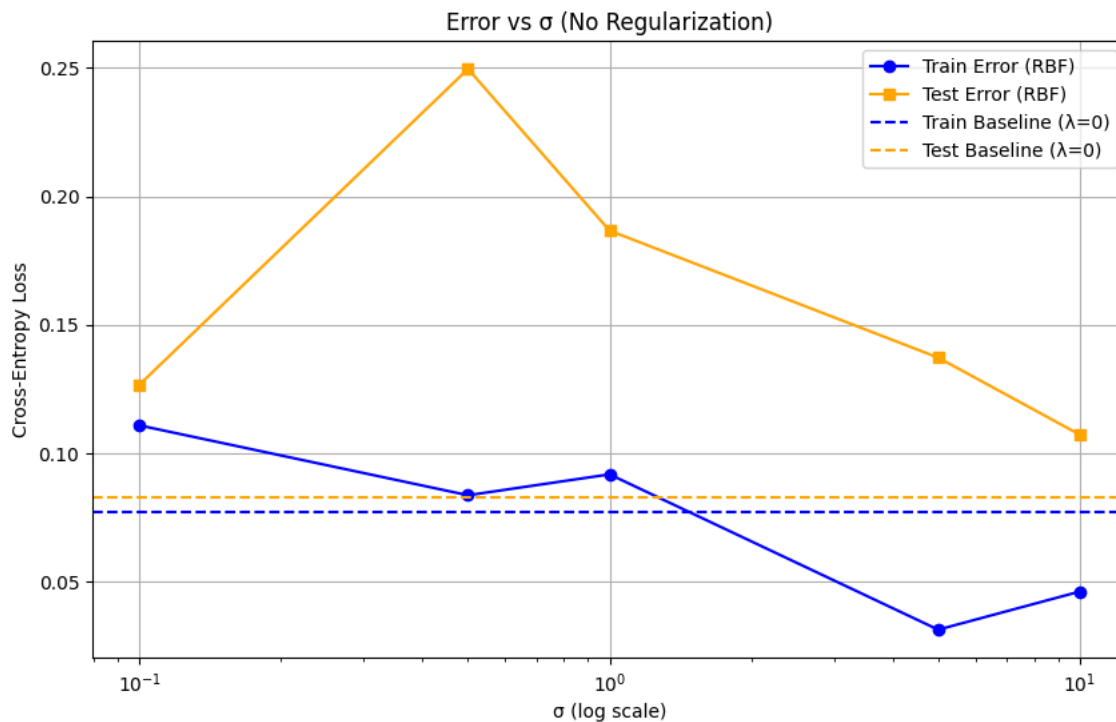
↔ Original feature dim: 30
   RBF-expanded feature dim: 750

```

```

1 #1(f)
2 # Description: For each  $\sigma$ , train logistic regression on only the basis functions for that  $\sigma$ .
3 # Plot train/test error vs  $\sigma$ . Also plot baseline from part (d) ( $\lambda=0$ ).
4
5 from sklearn.linear_model import LogisticRegression
6
7 # Store errors per sigma
8 train_errors_per_sigma = []
9 test_errors_per_sigma = []
10
11 # Also get baseline from part (d) with  $\lambda=0$  (first element)
12 baseline_train_loss = train_losses[0] #  $\lambda=0$ 
13 baseline_test_loss = test_losses[0]
14
15 # For each sigma, extract only its 150 features (30 orig features * 5 centers)
16 n_centers = 5
17 n_orig_features = X.shape[1]
18 start_idx = 0
19
20 for i, sigma in enumerate(sigmas):
21     # Extract features for this sigma
22     end_idx = start_idx + n_orig_features * n_centers
23     X_train_sigma = X_train_val_rbf[:, start_idx:end_idx]
24     X_test_sigma = X_test_rbf[:, start_idx:end_idx]
25     start_idx = end_idx
26
27     # Train model with NO regularization
28     model = LogisticRegression(penalty=None, solver='lbfgs', max_iter=1000, random_state=42)
29     model.fit(X_train_sigma, y_train_val)
30
31     # Predict and compute log loss
32     y_train_proba = model.predict_proba(X_train_sigma)
33     y_test_proba = model.predict_proba(X_test_sigma)
34     train_loss = log_loss(y_train_val, y_train_proba)
35     test_loss = log_loss(y_test, y_test_proba)
36
37     train_errors_per_sigma.append(train_loss)
38     test_errors_per_sigma.append(test_loss)
39
40 # Plot
41 plt.figure(figsize=(10, 6))
42 plt.plot(sigmas, train_errors_per_sigma, 'o-', label='Train Error (RBF)', color='blue')
43 plt.plot(sigmas, test_errors_per_sigma, 's-', label='Test Error (RBF)', color='orange')
44 plt.axhline(baseline_train_loss, color='blue', linestyle='--', label='Train Baseline ( $\lambda=0$ )')
45 plt.axhline(baseline_test_loss, color='orange', linestyle='--', label='Test Baseline ( $\lambda=0$ )')
46 plt.xscale('log')
47 plt.xlabel('σ (log scale)')
48 plt.ylabel('Cross-Entropy Loss')
49 plt.title('Error vs σ (No Regularization)')
50 plt.legend()
51 plt.grid(True)
52 plt.show()

```



```


1 #1(g)
2 # Description: Now use ALL 750 RBF features. Vary  $\lambda \in \{0, 0.1, \dots, 10000\}$ .
3 # Plot:
4 # 1. Train/Test loss vs  $\lambda$ 
5 # 2. L2 norm of full weight vector
6 # 3. L2 norm per  $\sigma$ -group vs  $\lambda$  (5 lines)
7
8 lambdas_extended = [0, 0.1, 1, 10, 100, 1000, 10000]
9 C_extended = [1e10 if lam == 0 else 1.0/lam for lam in lambdas_extended]
10
11 train_losses_rbf = []
12 test_losses_rbf = []
13 weight_norms_rbf = []
14 weight_norms_per_sigma = [] # list of lists: [ [norms for  $\sigma_1$ ], [ $\sigma_2$ ], ... ]
15
16 # Precompute slice indices for each sigma group
17 n_centers = 5
18 n_orig = X.shape[1]
19 sigma_slices = []
20 start = 0
21 for sigma in sigmas:
22     end = start + n_orig * n_centers
23     sigma_slices.append((start, end))
24     start = end
25
26 for lam, C in zip(lambdas_extended, C_extended):
27     model = LogisticRegression(penalty='l2', C=C, solver='lbfgs', max_iter=2000, random_state=42)
28     model.fit(X_train_val_rbf, y_train_val)
29
30     # Losses
31     y_train_proba = model.predict_proba(X_train_val_rbf)
32     y_test_proba = model.predict_proba(X_test_rbf)
33     train_loss = log_loss(y_train_val, y_train_proba)
34     test_loss = log_loss(y_test, y_test_proba)
35
36     # Full weight norm
37     w_full = model.coef_[0]
38     full_norm = np.linalg.norm(w_full)
39
40     # Norm per sigma group
41     norms_per_sigma = []
42     for start_idx, end_idx in sigma_slices:
43         w_group = w_full[start_idx:end_idx]
44         group_norm = np.linalg.norm(w_group)
45         norms_per_sigma.append(group_norm)
46
47     # Store
48     train_losses_rbf.append(train_loss)
49     test_losses_rbf.append(test_loss)
50     weight_norms_rbf.append(full_norm)
51     weight_norms_per_sigma.append(norms_per_sigma)

```

```

52
53 weight_norms_per_sigma = np.array(weight_norms_per_sigma) # shape (7, 5)
54
55 # Plot 1: Losses
56 plt.figure(figsize=(15, 5))
57 plt.subplot(1, 3, 1)
58 plt.semilogx(lambdas_extended, train_losses_rbf, 'o-', label='Train')
59 plt.semilogx(lambdas_extended, test_losses_rbf, 's-', label='Test')
60 plt.xlabel('λ')
61 plt.ylabel('Cross-Entropy')
62 plt.title('Loss vs λ (All RBFs)')
63 plt.legend()
64 plt.grid(True)
65
66 # Plot 2: Full weight norm
67 plt.subplot(1, 3, 2)
68 plt.semilogx(lambdas_extended, weight_norms_rbf, 'o-', color='purple')
69 plt.xlabel('λ')
70 plt.ylabel('||w||2')
71 plt.title('Full Weight Norm vs λ')
72 plt.grid(True)
73
74 # Plot 3: Weight norm per σ group
75 plt.subplot(1, 3, 3)
76 for i, sigma in enumerate(sigmas):
77     plt.semilogx(lambdas_extended, weight_norms_per_sigma[:, i], 'o-', label=f'σ={sigma}')
78 plt.xlabel('λ')
79 plt.ylabel('||wσ||2')
80 plt.title('Weight Norm per σ Group')
81 plt.legend()
82 plt.grid(True)
83
84 plt.tight_layout()
85 plt.show()

```



#### 1(h) Capturing Input Relationships with Gaussian Basis Functions

To capture relationships (interactions) between input variables, you need **multivariate Gaussian basis functions**:

$$\phi(x) = \exp(-0.5 * (x - \mu)^T \Sigma^{-1} (x - \mu))$$

Where:

- $\mu$  is a center vector in full input space ( $\mathbb{R}^d$ )
- $\Sigma$  is a covariance matrix (can be diagonal or full)

Impact on Bias-Variance Trade-off: Lower Bias: Can model complex interactions  $\rightarrow$  better fit to true function. Higher Variance: Many more parameters  $\rightarrow$  higher risk of overfitting unless strongly regularized or few centers used. Requires more data and careful tuning.

Alternative: Use tensor products of univariate RBFs — still captures interactions but suffers from curse of dimensionality (exponential growth in # of basis functions). """"

#### 1(i) Learning Algorithm for $\mu$ and $w$ (Fixed $\sigma$ , L2 on $w$ )

Let:

- $\phi(x; \mu) = \exp(-\|x - \mu\|^2 / (2\sigma^2))$  [Univariate RBF for simplicity, but can extend]

- Model:  $f(x) = \sum_j w_j \phi_j(x; \mu_j)$
- Loss:  $J(w, \mu) = (1/N) \sum_n \text{CE}(y_n, f(x_n)) + (\lambda/2) \|w\|_2^2$

Algorithm (Iterative Coordinate Descent):

1. Initialize centers  $\mu_j$  (e.g., random subset of  $X$ , or k-means centroids).
2. Repeat until convergence: a. **Fix  $\mu$ , optimize  $w$** : - Solve L2-regularized logistic regression (convex  $\rightarrow$  global optimum via LBFGS). b. **Fix  $w$ , optimize  $\mu$** : - Compute gradient for each  $\mu_i$ :  $\partial J / \partial \mu_i = (1/N) \sum_n [\sigma(f(x_n)) - y_n] * w_i * \phi_i(x_n; \mu_i) * (x_n - \mu_i) / \sigma^2$  - Update:  $\mu_i \leftarrow \mu_i - \eta * \partial J / \partial \mu_i$  (Gradient Descent)
3. Regularization: Applied only to  $w$  to avoid collapsing centers.

1(j) Does the algorithm converge? Local or Global?

- **Converges?** Yes, if step size  $\eta$  is chosen appropriately (e.g., via line search) and loss decreases monotonically. In practice, it often converges to a stationary point.
- **Global Optimum?** No. The joint optimization problem over  $w$  and  $\mu$  is **non-convex**. The loss surface has many local minima.
- **Local Optimum?** Yes. Gradient descent finds a local minimum. The quality of the solution heavily depends on initialization (e.g., initializing  $\mu$  with k-means usually gives better results than random).

This is a common challenge in models with hidden units or basis function centers (like RBF networks, neural nets, GMMs).

```

1 #2 Load and Preprocess California Housing Data
2 # Description: Load California Housing dataset, preprocess (scale, handle outliers, split),
3 # then apply Linear Regression, Ridge, Lasso, and Logistic Regression (classification version).
4
5 from sklearn.datasets import fetch_california_housing
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import StandardScaler
8 from sklearn.linear_model import LinearRegression, Ridge, Lasso, LogisticRegression
9 from sklearn.metrics import mean_squared_error, accuracy_score, classification_report
10 import pandas as pd
11
12 # Load dataset
13 cali = fetch_california_housing()
14 X, y = cali.data, cali.target
15
16 print("Original California Housing Dataset loaded.")
17
18 # Convert to classification: median_house_value > median → 1, else 0
19 y_class = (y > np.median(y)).astype(int)
20 print(f"Converted to binary classification. Class balance: {np.bincount(y_class)}")
21
22 # Create DataFrame for easier handling
23 df = pd.DataFrame(X, columns=cali.feature_names)
24 df['MedHouseVal'] = y
25 df['HighValue'] = y_class
26
27 # Check for missing values
28 print("Missing values:\n", df.isnull().sum())
29
30 # Outlier removal (optional): cap at 1.5 IQR
31 Q1 = df['MedHouseVal'].quantile(0.25)
32 Q3 = df['MedHouseVal'].quantile(0.75)
33 IQR = Q3 - Q1
34 lower_bound = Q1 - 1.5 * IQR
35 upper_bound = Q3 + 1.5 * IQR
36
37 df_clean = df[(df['MedHouseVal'] >= lower_bound) & (df['MedHouseVal'] <= upper_bound)]
38 print(f"Removed outliers. New shape: {df_clean.shape}")
39
40 # Prepare X and y
41 X_clean = df_clean[cali.feature_names].values
42 y_clean_reg = df_clean['MedHouseVal'].values
43 y_clean_class = df_clean['HighValue'].values
44
45 # Train-test split
46 X_train, X_test, y_train_reg, y_test_reg = train_test_split(
47     X_clean, y_clean_reg, test_size=0.2, random_state=42
48 )
49 _, _, y_train_class, y_test_class = train_test_split(
50     X_clean, y_clean_class, test_size=0.2, random_state=42
51 )
52
53 # Scale features
54 scaler = StandardScaler()
55 X_train_scaled = scaler.fit_transform(X_train)
56 X_test_scaled = scaler.transform(X_test)
57

```

```
58 print("Data preprocessing complete: scaled, split, outliers handled.")
```

```
➞ Original California Housing Dataset loaded.
Converted to binary classification. Class balance: [10323 10317]
Missing values:
  MedInc      0
  HouseAge    0
  AveRooms    0
  AveBedrms   0
  Population  0
  AveOccup    0
  Latitude    0
  Longitude   0
  MedHouseVal 0
  HighValue   0
dtype: int64
Removed outliers. New shape: (19569, 10)
Data preprocessing complete: scaled, split, outliers handled.
```

```
1 #Linear Regression and Regularized Versions (Regression Task)
2 # Description: Train Linear, Ridge, Lasso regression. Compare MSE.
3
4 models = {
5     'Linear': LinearRegression(),
6     'Ridge (α=1)': Ridge(alpha=1),
7     'Lasso (α=0.1)': Lasso(alpha=0.1, max_iter=10000)
8 }
9
10 print("=== REGRESSION RESULTS ===")
11 for name, model in models.items():
12     model.fit(X_train_scaled, y_train_reg)
13     y_pred = model.predict(X_test_scaled)
14     mse = mean_squared_error(y_test_reg, y_pred)
15     print(f"{name:15} Test MSE: {mse:.4f}")
```

```
➞ === REGRESSION RESULTS ===
Linear          Test MSE: 0.3688
Ridge (α=1)     Test MSE: 0.3688
Lasso (α=0.1)   Test MSE: 0.5166
```

```
1 #Logistic Regression (Classification Task)
2 # Description: Train logistic regression for classification (High/Low house value)
3
4 print("\n=== CLASSIFICATION RESULTS ===")
5 log_reg = LogisticRegression(max_iter=1000, random_state=42)
6 log_reg.fit(X_train_scaled, y_train_class)
7 y_pred_class = log_reg.predict(X_test_scaled)
8
9 acc = accuracy_score(y_test_class, y_pred_class)
10 print(f"Logistic Regression Accuracy: {acc:.4f}")
11 print("\nClassification Report:")
12 print(classification_report(y_test_class, y_pred_class, target_names=['Low', 'High']))
```

```
➞ === CLASSIFICATION RESULTS ===
Logistic Regression Accuracy: 0.8311
```


```
Classification Report:
              precision    recall  f1-score   support

     Low       0.85         0.83         0.84         2083
     High       0.81         0.83         0.82         1831

 accuracy                   0.83         0.83         0.83         3914
 macro avg              0.83         0.83         0.83         3914
 weighted avg           0.83         0.83         0.83         3914
```

```
1
2 # Plot coefficients of linear models
3 plt.figure(figsize=(12, 4))
4
5 for i, (name, model) in enumerate(models.items()):
6     if hasattr(model, 'coef_'):
7         plt.subplot(1, 3, i+1)
8         plt.bar(range(len(model.coef_)), model.coef_)
9         plt.title(name)
10        plt.xticks(range(len(cali.feature_names)), cali.feature_names, rotation=45, ha='right')
11
12 plt.tight_layout()
13 plt.show()
```





1 Start coding or [generate](#) with AI.